

Learning Deep Learning with PyTorch

(3) Knowing PyTorch

Qiyang Hu
IDRE

What is **PYTORCH**

- An open-source Python-based deep learning framework
 - Replacement for Numpy with supporting GPUs
 - A full set of deep learning libraries
- History
 - Lua-based Torch (2002 - 2011)
 - PyTorch 0.1 (2016): THNN
 - PyTorch 1.0 (2018): merging Caffe2
 - PyTorch 1.4 (Jan 15, 2020)
- About Fast.ai
 - Sitting on top of PyTorch
 - Fast.ai for PyTorch is not what Keras is for TF
 - Still unclear for the usage popularity

Why **PYTORCH**

- Simplicity & Flexibility from “Py”

- Feels like Numpy
- GPU acceleration

- Immediate execution mode

- Defined by run
- Tape-based autograd
- Awesome debugging

- Hybrid front-end

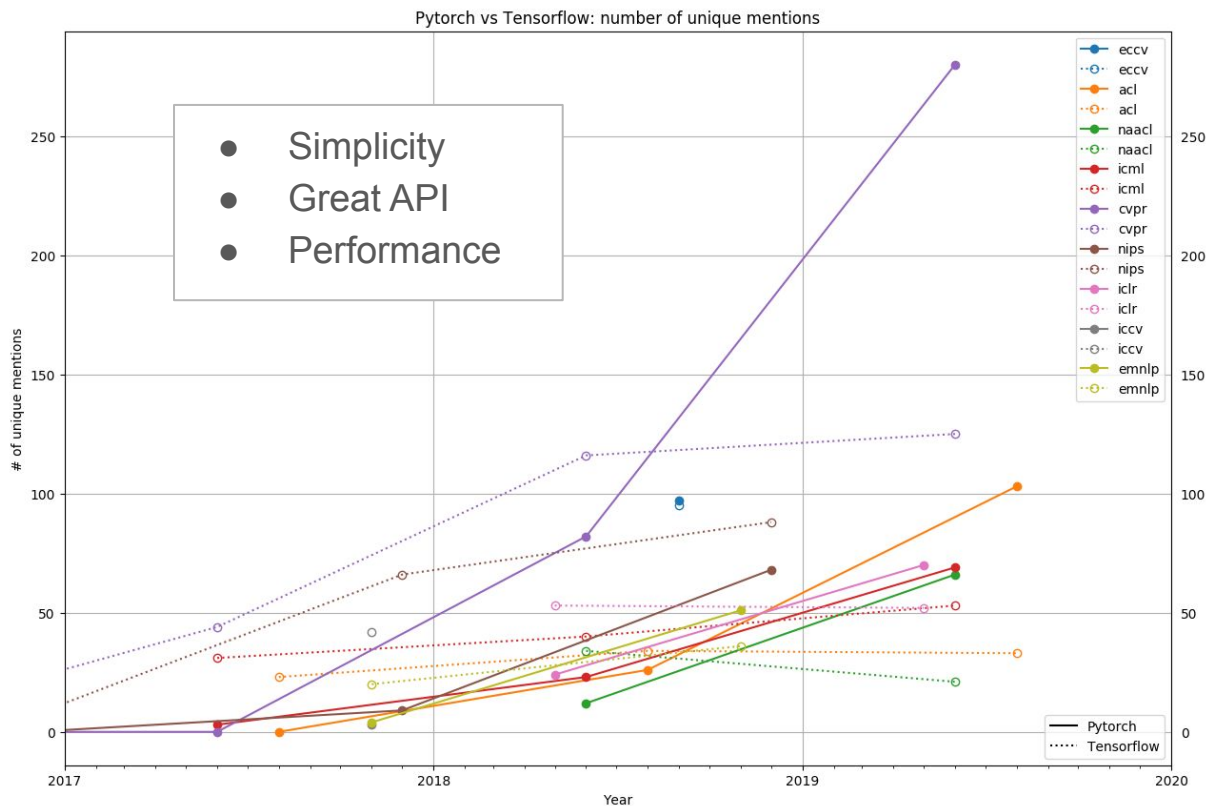
- JIT and TorchScript
- Seamlessly switch between
 - Modes
 - Distributed training
 - Mobile deployment

A graph is created on the fly

```
from torch.autograd import Variable  
  
x = Variable(torch.randn(1, 10))  
prev_h = Variable(torch.randn(1, 20))  
W_h = Variable(torch.randn(20, 20))  
W_x = Variable(torch.randn(20, 10))
```



PyTorch is increasing dominance in research



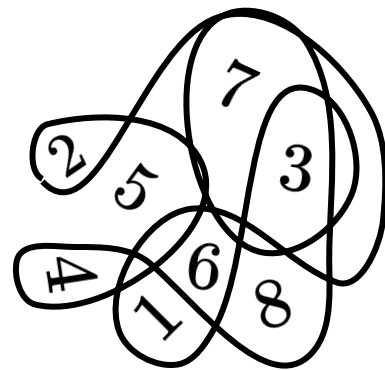
[Source](#)

Tensors as building blocks

1

$$\begin{bmatrix} 2 \\ 5 \\ 4 \end{bmatrix}$$

$$\begin{bmatrix} 2 & 1 & 3 \\ 5 & 7 & 9 \\ 4 & 8 & 6 \end{bmatrix}$$

$$\begin{bmatrix} \begin{bmatrix} 4 & 6 & 9 \\ 1 & 2 & 5 \\ 8 & 7 & 3 \end{bmatrix} \end{bmatrix}$$


Scalar
0-D

Vector
1-D

Matrix
2-D

Tensor
3-D

Tensor
 n -D

$$X = 1$$

$$X[1] = 5$$

$$X[2, 1] = 8$$

$$X[0, 1, 2] = 5$$

$$X[\underbrace{2, 3, \dots, 1}_{N \text{ indices}}] = 6$$

“Py” and “Non-Py” in PyTorch

- Tensors and Numpy arrays
 - Easy conversion
 - Zero copy: share their underlying memory locations (if on CPU)
 - Unboxed contiguous memory blocks containing unboxed C numeric types
 - Not python objects as in lists or tuples of numbers
 - Views by tensor metadata, e.g. sizes, strides, offsets
- PyTorch = Python + C/C++ + CUDA
 - Python extension objects in C/C++
 - Code base components:
 - The core Torch libraries: TH, THC, THNN, THCUNN
 - Vendor libraries: CuDNN, NCCL
 - Python Extension libraries
 - Additional third-party libraries: NumPy, MKL, LAPACK, DLPack

Colab Hands-on

bit.ly/LDDL_01

Automatic differentiation

- Autograd package
 - Track all operations of tensors
 - Compute derivatives analytically via back-prop
 - Natively loaded in torch module
 - Can be used in other scientific domains
- Simple usage
 - Set tensor's `.requires_grad` as `TRUE`
 - Call `.backward()`
 - Gradient accumulated into `.grad` attribute
 - Tensor's creation function recorded in `.grad_fn` attribute
- Stop a tensor from tracking history
 - `.detach()`
 - Wrap the code block in with `torch.no_grad()`

Neural Networks in PyTorch

- [torch.nn](#) package
 - Contains all building blocks for NN architectures
 - All blocks subclassed from `nn.Module` (e.g. `nn.Linear`)
- Define a network
 - For simple networks:
 - concatenate modules through the `nn.Sequential` container
 - For complex networks:
 - Subclassing `nn.Module`
- `nn.Module` package expects first index as first batch size of samples
 - Need to reshape the input by `.unsqueeze()`
- Loss functions in `torch.nn`:
 - `L1Loss`, `MSELoss`, `CrossEntropyLoss`, `MarginRankingLoss`, ...

Optimizers in PyTorch

- [torch.optim](#) package
 - Provides various optimization algorithms
 - Need to move model to GPU before constructing optimizers
 - Must zero the gradient explicitly:
 - `optimizer.zero_grad()`
 - Take an optimization step:
 - `optimizer.step()` in GD method
 - `optimizer.step(closure)` in CG or LBFGS method
 - Optional: adjust the learning rate based on the number of epochs.
 - `optimizer.lr_scheduler`

Don't forget to

- Sign in your info to the class
 - To get the email notifications
- Contact me for questions or discussions
 - huqy@idre.ucla.edu
 - Office: Math Sci #3330
 - Phone: 310-825-2011